

1A

1. Given an array of integers $X = \{x_0, \dots, x_{n-1}\}$, we want to partition the array into k groups P_1, \dots, P_k so that the numbers in every group P_j are distinct, minimizing k , and preserving order.

An algorithm is as follows. Start with array X . Create group $P = \{\}$. Let $k = 0$.

While X is not empty, remove the first x from X . If $x \notin P$, then add x to P . If x already in P , then let $k = k + 1$ and reset $P = \{\}$.

The solution to the problem is k when X is empty. (Note that while the partitions P are not unique, k is unique.)

2. Now we prove that the algorithm is correct.

Proof. Let k_A be the solution given by the algorithm, and let k_O be the optimal solution. We know that $k_A \not< k_O$, since that would contradict the optimality of k_O .

Seeking a contradiction, suppose that $k_O < k_A$, i.e. that there are strictly fewer partitions in the optimal solution than the number in the algorithm solution. This means that all elements in P_{k_A} (i.e. all elements in the last partition of the algorithm solution) belong to some partition P'_i in the optimal solution, where $i < k_A$.

This is a contradiction, because our algorithm only creates a new partition when adding another element to a partition would mean that that partition had duplicate elements. In particular, the first element in P_{k_A} (the first element in the last partition of the algorithm solution), cannot belong to P_{k_A-1} , since by the algorithm construction, that element's number is already in P_{k_A-1} , and it certainly cannot belong to any other partition, since then the elements would be out of order.

Thus, it cannot be that $k_O < k_A$, and we also know that $k_A \not< k_O$. We therefore conclude that $k_O = k_A$, i.e. the algorithm solution is optimal and our algorithm is correct. \square

3. Finally we analyze the running time. At the i th step of the algorithm, we need to decide if the i th element of X is contained in a partition P which has at most $i - 1$ elements. Finding an element in a list of size n can be implemented in $O(\log n)$ time. At each step we also do some $O(1)$ operations. We have a step for each element in X , which has length n . Thus the total running time of the algorithm is $O(n \log n)$.

1B

1. Given an array of integers $X = \{x_0, \dots, x_{n-1}\}$, we want to partition the array into k groups P_1, \dots, P_k so that the numbers in every group P_j are distinct, minimizing k , but we *don't* care about preserving order.

An algorithm is as follows:

Create a map M , of elements (x, c) where $x \in X$, and c is the number of times x appears in X .

While X not empty, remove the first element of X . If x is not a key in M , then add $(x, 1)$ to M . If x is a key in M , i.e. $(x, c) \in M$ for some c , then increment c so that $(x, c + 1) \in M$.

When X is empty, find the highest value c in M . This is the minimum number of partitions necessary.

(In other words, find the element with the most number of repetitions in X , the number of repetitions is k .)

2. Now we prove that the algorithm is correct.

Proof. Let k_A be the number of partitions returned by the algorithm, and let k_O be the optimal (minimum) number of partitions. Clearly, $k_A \not< k_O$, since that would contradict the optimality of k_O . Seeking a contradiction, suppose $k_O < k_A$, i.e. that there are strictly fewer partitions in the optimal solution than the number in the algorithm solution. However, by the construction of the algorithm, we know that there must be some element x which appears in X k_A times. But $k_O < k_A$, so if x appears k_A times then there must be some partition in the optimal solution which contains a repetition of an element, which is not allowed. Therefore it cannot be that $k_O < k_A$. Thus if $k_O \not< k_A$ and $k_A \not< k_O$, then we must have $k_A = k_O$. Our algorithm is therefore correct. \square

3. Finally we analyze the running time. We have a step for each element in X , which has length n . With appropriate implementation, we can find/update a particular element in the map M in $O(\log n)$ time (for `c++`, we can use the `map` object). Finding the element with the highest number of repetitions takes $O(n)$. Together the running time is $n \times O(\log n) + O(n) = O(n \log n)$.